

# WEEK 2

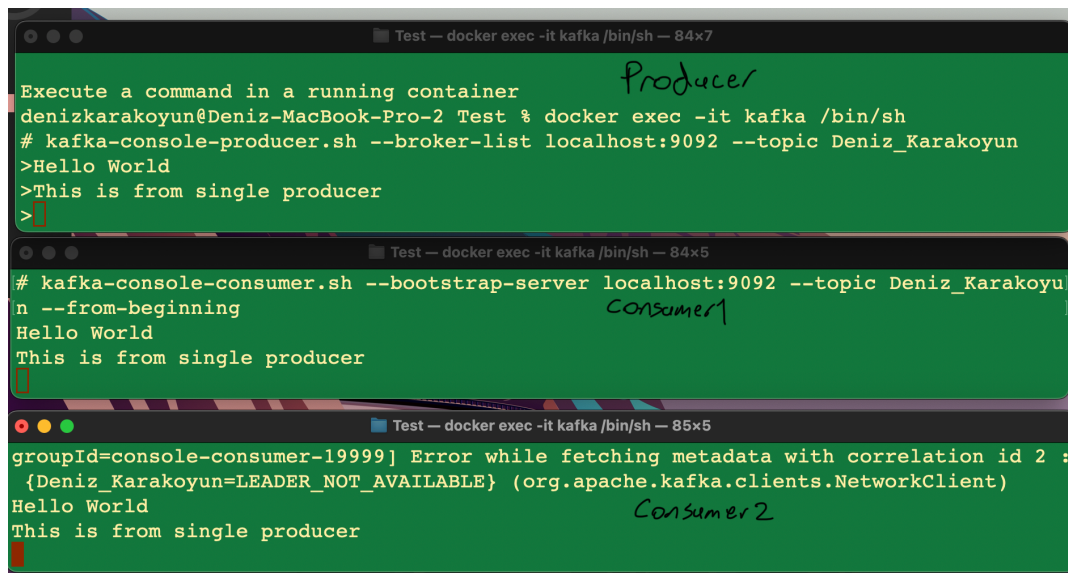
August 23, 2024

## Student Information

Full Name: Deniz Karakoyun

Id Number: 2580678

### 1) Multiple Consumer Single Producer



The image shows three terminal windows stacked vertically, all with a green background. The top window is titled 'Test -- docker exec -it kafka /bin/sh -- 84x7' and contains the following text: 'Execute a command in a running container' (handwritten), 'denizkarakoyun@Deniz-MacBook-Pro-2 Test % docker exec -it kafka /bin/sh' (handwritten), '# kafka-console-producer.sh --broker-list localhost:9092 --topic Deniz\_Karakoyun' (handwritten), '>Hello World' (handwritten), '>This is from single producer' (handwritten), and '>' (handwritten). The middle window is titled 'Test -- docker exec -it kafka /bin/sh -- 84x5' and contains: '# kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic Deniz\_Karakoyun' (handwritten), 'n --from-beginning' (handwritten), 'Hello World' (handwritten), and 'This is from single producer' (handwritten). The bottom window is titled 'Test -- docker exec -it kafka /bin/sh -- 85x5' and contains: 'groupId=console-consumer-19999] Error while fetching metadata with correlation id 2 : {Deniz\_Karakoyun=LEADER\_NOT\_AVAILABLE} (org.apache.kafka.clients.NetworkClient)' (handwritten), 'Hello World' (handwritten), and 'This is from single producer' (handwritten). Handwritten labels 'Producer' and 'Consumer1' are next to the producer and first consumer windows respectively, and 'Consumer2' is next to the second consumer window.

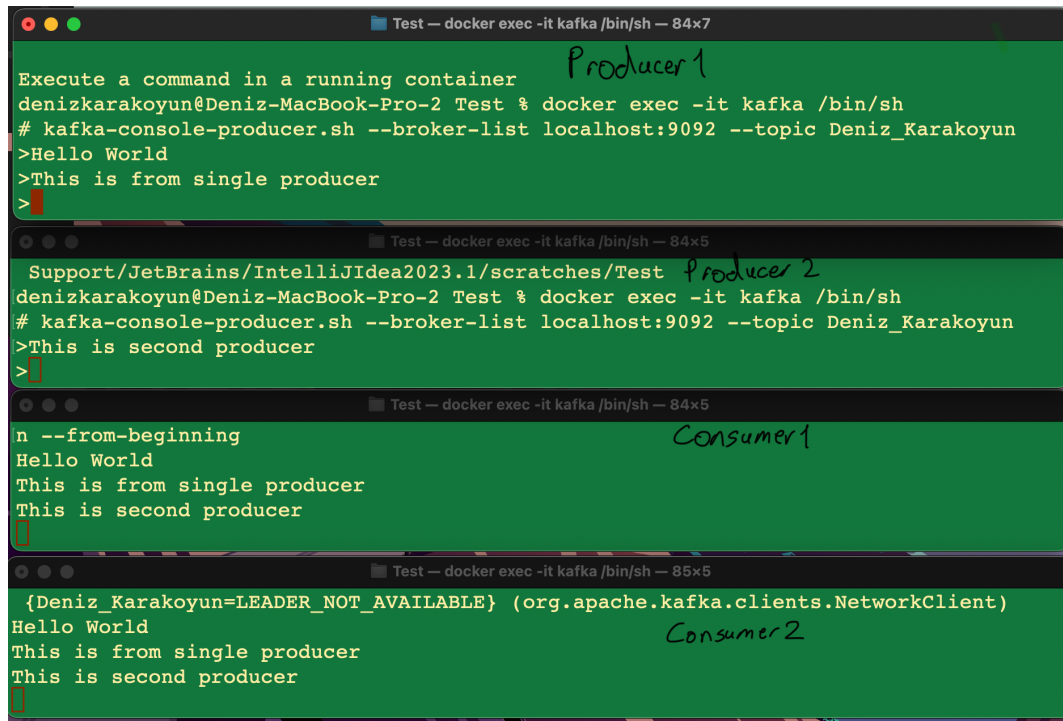
Figure 1: Multiple Consumer Single Producer Architecture

#### Explanation:

In this architecture, a single Kafka producer sends messages to a Kafka topic, which can be consumed by multiple consumers. This setup is useful for distributing the processing load among multiple consumers. Each consumer can read from the same topic independently, ensuring that

the processing of messages is parallelized. This approach is often used to scale out the processing capabilities of an application.

## 2) Multiple Consumer Multiple Producer



The image displays four terminal windows stacked vertically, each running a Kafka command-line interface. The first window, titled 'Test - docker exec -it kafka /bin/sh - 84x7', shows a producer sending messages: 'Hello World' and 'This is from single producer'. The second window, titled 'Test - docker exec -it kafka /bin/sh - 84x5', shows another producer sending 'This is second producer'. The third window, titled 'Test - docker exec -it kafka /bin/sh - 84x5', shows a consumer receiving all three messages. The fourth window, titled 'Test - docker exec -it kafka /bin/sh - 85x5', shows a second consumer receiving the same three messages, with a warning message '{Deniz\_Karakoyun=LEADER\_NOT\_AVAILABLE} (org.apache.kafka.clients.NetworkClient)' appearing before the first message.

```
Test - docker exec -it kafka /bin/sh - 84x7
Execute a command in a running container
denizkarakoyun@Deniz-MacBook-Pro-2 Test % docker exec -it kafka /bin/sh
# kafka-console-producer.sh --broker-list localhost:9092 --topic Deniz_Karakoyun
>Hello World
>This is from single producer
>

Test - docker exec -it kafka /bin/sh - 84x5
Support/JetBrains/IntelliJ IDEA 2023.1/scratches/Test
denizkarakoyun@Deniz-MacBook-Pro-2 Test % docker exec -it kafka /bin/sh
# kafka-console-producer.sh --broker-list localhost:9092 --topic Deniz_Karakoyun
>This is second producer
>

Test - docker exec -it kafka /bin/sh - 84x5
n --from-beginning
Hello World
This is from single producer
This is second producer

Test - docker exec -it kafka /bin/sh - 85x5
{Deniz_Karakoyun=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)
Hello World
This is from single producer
This is second producer
```

Figure 2: Multiple Consumer Multiple Producer Architecture


### Explanation:

This more complex architecture involves multiple producers sending messages to Kafka topics, which are then consumed by multiple consumers. This setup offers high flexibility and fault tolerance. Multiple producers can independently send data to the same or different topics, and multiple consumers can read from these topics. This model is particularly useful in scenarios requiring high availability and scalability, as both producers and consumers can operate independently and handle failures gracefully.

## 3) Multiple Partitions

Kafka topics are divided into partitions to allow parallel processing and increase performance. Each partition is an ordered log of messages. Understanding how partitions work is crucial for optimizing Kafka performance and ensuring data is processed efficiently.

## Creating a Topic with Multiple Partitions




```
Test — docker exec -it kafka /bin/sh — 84x7
#
# kafka-topics.sh --create --zookeeper zookeeper:2181 --replication-factor 1 --partitions 3 --topic multiple_partition
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.
Created topic multiple_partition.
#
```

Figure 3: Creating a Topic with Multiple Partitions

### Explanation:

Creating a topic with multiple partitions allows Kafka to distribute the load of incoming messages across multiple brokers. This setup enhances performance by enabling parallel processing. Each partition can be stored on a different broker, and producers write messages to partitions based on a configured partitioning strategy. This approach also improves fault tolerance, as the failure of one partition does not affect the others.

## Viewing Partition Logs




> 📁 multiple_partition-0	VOLUME	4 minutes ago	drwxr-xr-x
> 📁 multiple_partition-1	VOLUME	4 minutes ago	drwxr-xr-x
> 📁 multiple_partition-2	VOLUME	4 minutes ago	drwxr-xr-x

Figure 4: Logs of Multiple Partitions

### Explanation:

The logs of multiple partitions display the contents and status of each partition. Kafka maintains separate logs for each partition, and these logs are managed independently. Understanding these logs helps in monitoring the data flow and diagnosing issues related to message distribution and partitioning.

## Writing Messages to a Topic with Multiple Partitions



```
Test — docker exec -it kafka /bin/sh — 84x7
denizkarakoyun@Deniz-MacBook-Pro-2 Test % docker exec -it kafka /bin/sh
# kafka-console-producer.sh --topic multiple_partition --bootstrap-server localhost:9092
>Message 1
>Message 2
>Message 3
>^C#
```

Figure 5: Writing Messages to a Topic with Multiple Partitions

**Explanation:**

When messages are written to a topic with multiple partitions, Kafka distributes these messages based on the configured partitioning strategy. Typically, messages are assigned to partitions using a key-based or round-robin approach. This distribution ensures that the workload is evenly spread and allows consumers to process messages from different partitions in parallel.

**Partition Content Analysis**

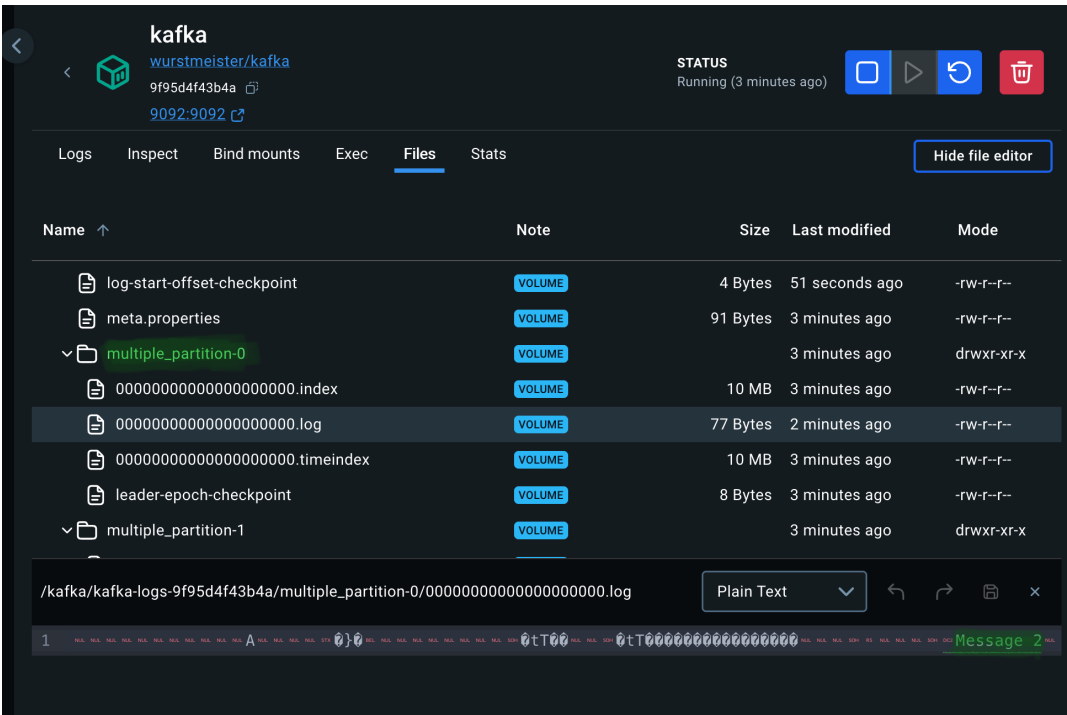


Figure 6: Partition 0 and Its Content

**Explanation:**

Partition 0 contains the message "Message 2". Each partition holds an ordered log of messages. Understanding the content of each partition helps in analyzing the message distribution and ensures that data is processed as expected.

kafka

wurstmeister/kafka

9f95d4f43b4a

9092:9092

STATUS

Running (4 minutes ago)

Logs

Inspect

Bind mounts

Exec

Files

Stats

Hide file editor

Name ↑	Note	Size	Last modified	Mode
00000000000000000000.timeindex	VOLUME	10 MB	4 minutes ago	-rw-r--r--
leader-epoch-checkpoint	VOLUME	8 Bytes	4 minutes ago	-rw-r--r--
multiple_partition-1	VOLUME		4 minutes ago	drwxr-xr-x
00000000000000000000.index	VOLUME	10 MB	4 minutes ago	-rw-r--r--
00000000000000000000.log	VOLUME	0 Bytes	4 minutes ago	-rw-r--r--
00000000000000000000.timeindex	VOLUME	10 MB	4 minutes ago	-rw-r--r--
leader-epoch-checkpoint	VOLUME	8 Bytes	4 minutes ago	-rw-r--r--
multiple_partition-2	VOLUME		4 minutes ago	drwxr-xr-x
00000000000000000000.index	VOLUME	10 MB	4 minutes ago	-rw-r--r--

/kafka/kafka-logs-9f95d4f43b4a/multiple\_partition-1/00000000000000000000000000000000.log

Plain Text

1

Figure 7: Partition 1 and Its Content

### Explanation:

Partition 1 currently has no messages. This could be due to a variety of reasons, such as the partition not receiving any data or messages being distributed to other partitions based on the partitioning strategy.

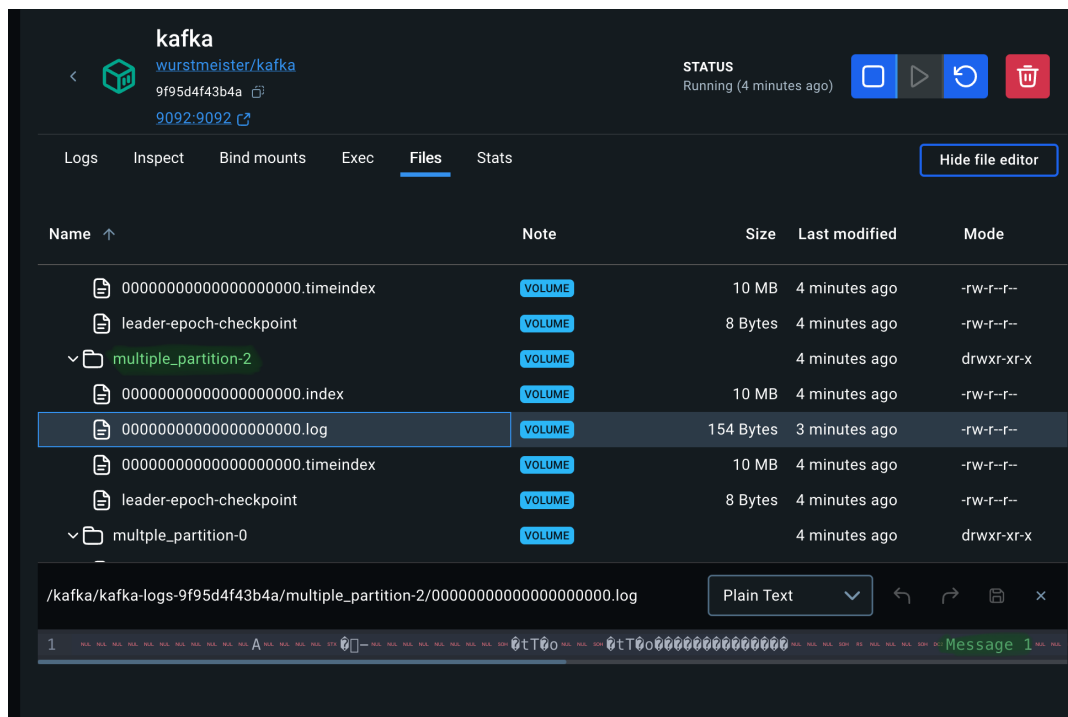


Figure 8: Partition 2 and Its Content

**Explanation:**

Partition 2 contains the messages "Message 1" and "Message 3". Each partition can hold multiple messages, and the order of messages is preserved within a partition. Analyzing the contents of each partition helps in understanding the message distribution and verifying the correctness of the data flow.

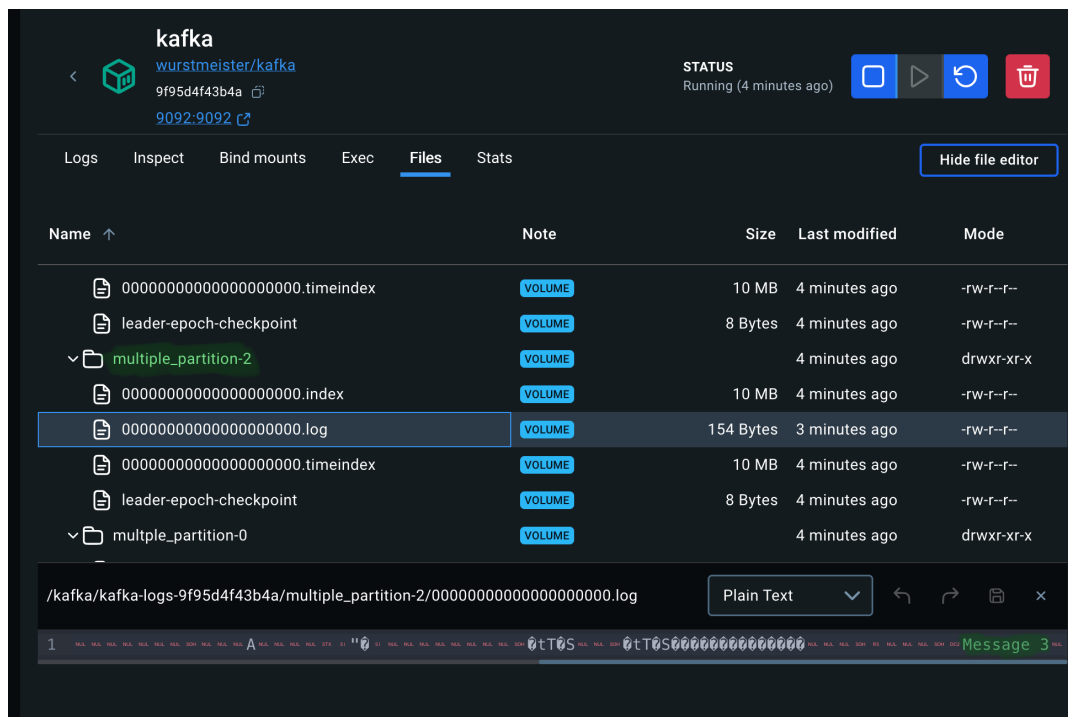


Figure 9: Detailed View of Partition 2

**Explanation:**

This detailed view of Partition 2 shows the messages "Message 1" and "Message 3" in more detail. Such detailed analysis can be useful for debugging and ensuring that messages are being distributed and stored correctly within each partition.

#### 4) Multiple Partitions: Reading a Topic

In Kafka, topics are divided into multiple partitions, facilitating parallel processing and enhanced performance. Understanding how to consume messages from these partitions is crucial for effective data management and analysis.

## Consuming Messages from a Topic with Multiple Partitions

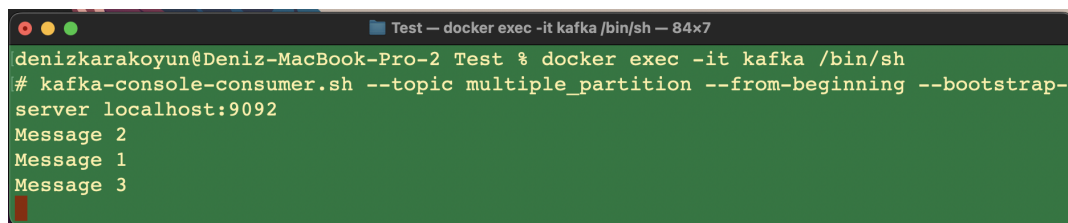


Figure 10: Consuming Messages From a Topic with Multiple Partitions

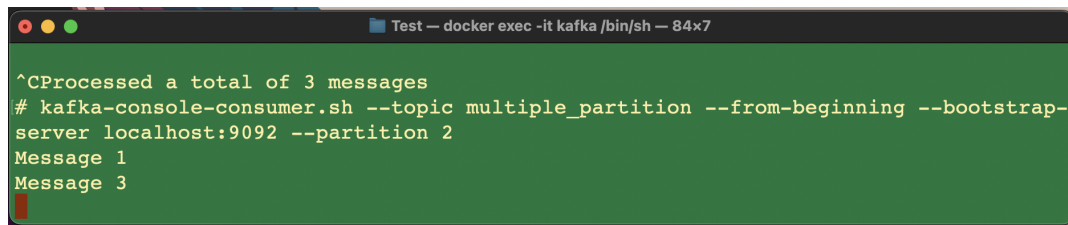
### Explanation:

When consuming messages from a topic with multiple partitions, Kafka allows parallel reading of messages from all partitions. This parallelism improves the efficiency and scalability of data processing.

- **Consumers:** A consumer (or consumer group) subscribes to a topic and reads messages from all partitions associated with that topic.
- **Partition Assignment:** Each consumer in the group is assigned one or more partitions. Kafka ensures that each partition is only read by one consumer within a group at any given time.
- **Message Distribution:** Messages are distributed among partitions based on the producer's partitioning strategy. Consumers then read messages from their assigned partitions.

The figure above demonstrates how multiple consumers can read messages from different partitions of the same topic, enhancing parallel processing and system scalability.

## Consuming Messages from a Specified Partition

A terminal window titled "Test — docker exec -it kafka /bin/sh — 84x7" with a green background. It shows the output of a Kafka console consumer command. The output indicates that 3 messages were processed, and the consumer is reading from partition 2, displaying "Message 1" and "Message 3".

```
^CProcessed a total of 3 messages
# kafka-console-consumer.sh --topic multiple_partition --from-beginning --bootstrap-server localhost:9092 --partition 2
Message 1
Message 3
```

Figure 11: Consuming Messages From a Specified Partition

### Explanation:

Sometimes, it is necessary to consume messages from a specific partition rather than from all partitions of a topic. This targeted approach is useful for tasks such as debugging or analyzing specific data subsets.

- **Partition Specification:** Configure your consumer to read from a specific partition by specifying its number in the consumer configuration.
- **Consumer Configuration:** Programmatically create a `TopicPartition` object and assign it to the consumer to interact with the desired partition.
- **Use Case:** This method is beneficial for isolating and analyzing messages from a particular partition or when troubleshooting distribution issues.

The figure illustrates how a consumer can be configured to read from a specified partition, focusing on targeted message retrieval.

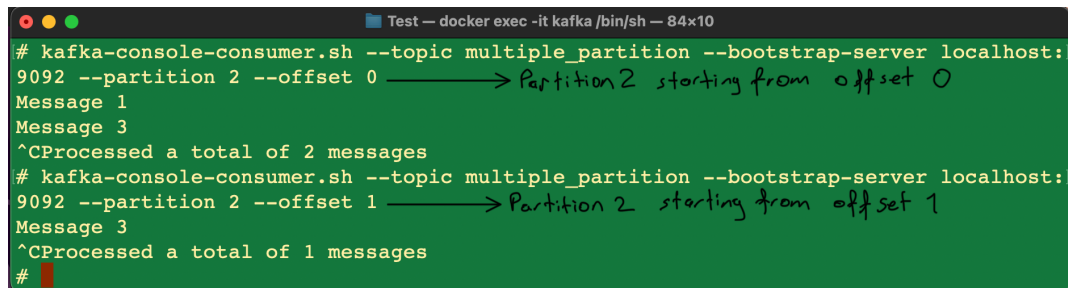


## Consuming Messages from a Specified Partition and Offset

### Explanation:

In addition to specifying a partition, you may want to start consuming messages from a specific offset within that partition. This functionality is important for processing messages from a specific point or event.

- **Offset Specification:** Set the offset from which to start consuming messages. This is useful for resuming from a known point or skipping certain messages.
- **Consumer Configuration:** Use the `assign` method to set both the partition and offset directly, or use the `seek` method to move the consumer's position to the desired offset.
- **Use Case:** Helpful for message replay, failure recovery, or analyzing specific data ranges.



```
Test — docker exec -it kafka /bin/sh — 84x10
# kafka-console-consumer.sh --topic multiple_partition --bootstrap-server localhost:
9092 --partition 2 --offset 0 —————> Partition 2 starting from offset 0
Message 1
Message 3
^CProcessed a total of 2 messages
# kafka-console-consumer.sh --topic multiple_partition --bootstrap-server localhost:
9092 --partition 2 --offset 1 —————> Partition 2 starting from offset 1
Message 3
^CProcessed a total of 1 messages
#
```

Figure 12: Consuming Messages From a Specified Partition and Specified Offset

The consumer will begin reading from the specified offset within the partition, providing precise control over message consumption.

### Summary:

- **General Process:** Consuming messages from all partitions of a topic enhances parallel processing.
- **Specified Partition:** Allows targeted message retrieval from a single partition.
- **Specified Partition and Offset:** Provides fine-grained control over message consumption starting from a specific point.

These are the parts of the process of consuming messages from Kafka topics with multiple partitions, as well as handling specific partitions and offsets. Adjust the figures and paths as necessary to align with your specific setup and use case.

## 5) Multiple Brokers

In Apache Kafka, a **broker** is a server responsible for receiving, storing, and serving messages. Kafka's architecture is designed to be **distributed**, meaning it can run on multiple brokers, pro-

viding scalability, fault tolerance, and increased performance.

## Key Concepts

- **Broker:** A single Kafka server that handles producer requests, stores data, and serves consumers.
- **Cluster:** A group of Kafka brokers working together. In a Kafka cluster, data is distributed across multiple brokers.
- **Partitioning:** Each topic in Kafka is divided into multiple partitions, which are distributed across different brokers. This enables parallel processing and better load distribution.
- **Replication:** Kafka replicates each partition across multiple brokers, ensuring data durability and availability even in the event of broker failures.
- **Leader and Follower:** For each partition, one broker acts as the leader, and the others are followers. The leader handles all read and write requests for the partition, while followers replicate the data. If a leader broker fails, a follower is automatically elected as the new leader.
- **High Availability:** With multiple brokers, Kafka can continue operating even if one or more brokers fail, as data is replicated across the cluster.

By distributing data across multiple brokers, Kafka can achieve high throughput, fault tolerance, and scalability, making it suitable for large-scale, real-time data streaming applications.

```

denizkarakoyun@Deniz-MacBook-Pro-2 Test % docker compose -f docker-compose-1.yml up -d
WARN[0000] /Users/denizkarakoyun/Library/Application Support/JetBrains/IntelliJ IDEA 2023.1/scratches/Test/docker-compose-1.yml: the
attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion
[+] Running 17/17
  ✓ kafka2 Pulled                    56.7s
  ✓ kafka3 Pulled                    56.7s
  ✓ 96fb404a82ee Pull complete       14.9s
  ✓ b898cebac917 Pull complete       14.9s
  ✓ 09d715bc4459 Pull complete       51.9s
  ✓ 255391c00c77 Pull complete       51.9s
  ✓ f2224cc31197 Pull complete       51.9s
  ✓ e76d01693e7f Pull complete      52.1s
  ✓ 4e7863a5f3c1 Pull complete      52.1s
  ✓ 8643701ade91 Pull complete      52.1s
  ✓ 569add7486b2 Pull complete      52.1s
  ✓ e8d6727a787c Pull complete      52.7s
  ✓ c412fcd16cb4 Pull complete      52.7s
  ✓ zookeeper Pulled                56.7s
  ✓ ec3cd468da30 Pull complete       52.7s
  ✓ da70cf005879 Pull complete       52.7s
  ✓ kafka1 Pulled                    56.7s
WARN[0056] Found orphan containers ([kafka]) for this project. If you removed or renamed this service in your compose file, you ca
n run this command with the --remove-orphans flag to clean it up.
[+] Running 4/4
  ✓ Container zookeeper Started      0.8s
  ✓ Container kafka2 Started         0.5s
  ✓ Container kafka1 Started         0.5s
  ✓ Container kafka3 Started         0.5s
denizkarakoyun@Deniz-MacBook-Pro-2 Test % docker ps

```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
260def9e3e18	confluentinc/cp-kafka:7.2.1	"/etc/confluent/dock..."	5 minutes ago	Up 5 minutes	0.0.0.0:8098->8098/tcp, 9
092/tcp	kafka2				
f8bc909f05ed	confluentinc/cp-kafka:7.2.1	"/etc/confluent/dock..."	5 minutes ago	Up 5 minutes	0.0.0.0:8097->8097/tcp, 9
092/tcp	kafka1				
d1d269878e9b	confluentinc/cp-kafka:7.2.1	"/etc/confluent/dock..."	5 minutes ago	Up 5 minutes	0.0.0.0:8099->8099/tcp, 9
092/tcp	kafka3				
8422b9e23a8d	confluentinc/cp-zookeeper:7.2.1	"/etc/confluent/dock..."	5 minutes ago	Up 5 minutes	2181/tcp, 2888/tcp, 3888/
tcp	zookeeper				

```

denizkarakoyun@Deniz-MacBook-Pro-2 Test %

```

Figure 13: This figure illustrates the configuration of multiple brokers in Kafka by setting server properties. Each broker has a unique configuration, allowing them to work together in a cluster.

```

1  docker-compose-1.yml
2  services:
3    zookeeper:
4      image: confluentinc/cp-zookeeper:7.2.1
5      container_name: zookeeper
6      environment:
7        ZOOKEEPER_CLIENT_PORT: 2181
8    kafka1:
9      image: confluentinc/cp-kafka:7.2.1
10     container_name: kafka1
11     ports:
12       - "8097:8097"
13     depends_on:
14       - zookeeper
15     environment:
16       KAFKA_BROKER_ID: 1
17       KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
18       KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: EXTERNAL:PLAINTEXT,INTERNAL:PLAINTEXT
19       KAFKA_ADVERTISED_LISTENERS: EXTERNAL://localhost:8097,INTERNAL://kafka1:9092
20       KAFKA_INTER_BROKER_LISTENER_NAME: INTERNAL
21    kafka2:
22      image: confluentinc/cp-kafka:7.2.1
23      container_name: kafka2
24      ports:
25        - "8098:8098"
26      depends_on:
27        - zookeeper
28      environment:
29        KAFKA_BROKER_ID: 2
30        KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
31        KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: EXTERNAL:PLAINTEXT,INTERNAL:PLAINTEXT
32        KAFKA_ADVERTISED_LISTENERS: EXTERNAL://localhost:8098,INTERNAL://kafka2:9092
33        KAFKA_INTER_BROKER_LISTENER_NAME: INTERNAL
34    kafka3:
35      image: confluentinc/cp-kafka:7.2.1
36      container_name: kafka3
37      ports:
38        - "8099:8099"
39      depends_on:
40        - zookeeper
41      environment:
42        KAFKA_BROKER_ID: 3
43        KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
44        KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: EXTERNAL:PLAINTEXT,INTERNAL:PLAINTEXT
45        KAFKA_ADVERTISED_LISTENERS: EXTERNAL://localhost:8099,INTERNAL://kafka3:9092

```

Figure 14: This figure shows the regulation and tuning of server settings to optimize performance across multiple brokers. Proper regulation ensures efficient resource utilization and fault tolerance.

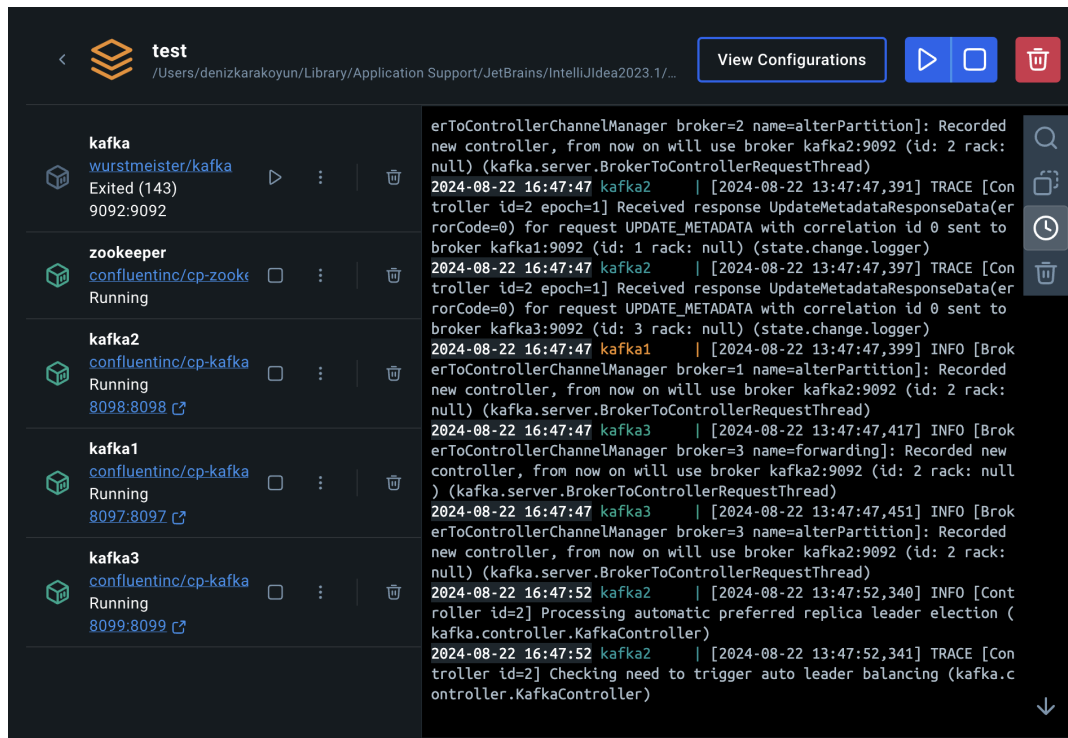


Figure 15: This figure demonstrates the process of starting multiple brokers in a Kafka cluster. Each broker is initialized, and they communicate to form a cohesive cluster, enabling distributed data processing.